

Backpropagation

JingYang Zeng, David Pomerence

Advanced Concepts of Machine Learning (Kurt Driessens) 2020

Department of Data Science and Knowledge Engineering, Maastricht University

Abstract—We implement the backpropagation algorithm for artificial neural networks, and applied it to a simple synthetic classification problem. Our implementation achieves an optimal encoding with a modest training time of only 50 epochs.

I. IMPLEMENTATION

We implement the neural network in Python, making use of the *numpy* library¹ for fast vector mathematics. We mainly follow the implementation by Ng[1], enriching it with additional alternative activation functions.

A. Documentation

The core functionality is found in *src/neural_network.py*.

The training data (both input and output) is of the form of vectors of a fixed size n_l , where exactly one value is set to 1, and all other values are set to 0. These data are generated by the function *generate_dataset*, which takes the size as an argument and returns the list of all possible data with this size.

We store the weights W_l and b_l in a dictionary called *parameters*. Initial random values for these parameters are created with *initialize_parameters*, optionally taking a seed number as an argument.

Once the test data are generated and the parameters are initialized, a network can be trained with *train_model*. This takes the input data, output data, an array indicating the shape of the neural network, and the two hyperparameters: learning rate and iteration count. This function can create networks of any size and shape, so it is more general than required for this assignment; but we will make some use of this advantage in our analysis of the weights.

The functionality related to the activation functions is provided in *src/kernel.py*. We provide implementations of multiple activation functions:

The sigmoid function $f(x) = \frac{1}{1+e^{-x}}$, the hyperbolic tangent $f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$, the softmax function $f(x) = \frac{e^x}{\sum_{i=1}^n e^{x_i}}$, the rectifier $f(x) = \max(0, x)$, and the linear function $f(x) = x$, and their respective derivatives.

forward_propagation performs a forward propagation from the activations of one level of units, returning the activations of the next level, taking the activation function to be used as an argument. *backward_propagation* does the same for the backpropagation of the errors. We do not use a regularization term for the calculation of the errors, thus there is no weight decay parameter.

tests/test.py contains the code for the analysis further described in this document.

¹<https://numpy.org>

B. Instructions

We use *Poetry*²³ as a deterministic project dependency manager. The code can be executed as follows:

```
poetry install
poetry run python tests/test.py
```

Alternatively, we have provided a *requirements.txt* for usage with *pip*. Since—unlike *poetry-pip* installs dependencies globally, this will require the manual setup of a virtual environment.

II. LEARNING PERFORMANCE

We encountered the following problem: When we implement the neural network, at first we mistakenly used a sigmoid function for our last hidden layer; with that configuration, even after iterating 10000 times there was still an overall cost of 0.77, and we could only achieve an accuracy up to 0.625.

Then we realized our task is a multiple classification rather than binary classification; the sigmoid function would like a Bernoulli distribution, each value in result is independent of all other values, but the expected results are 8 distinguished predictions and should only have a “1” in each result. Therefore, we used *softmax* and cross-entropy loss, which works for multiple classification, as an activation function for our last layer, which dramatically improved the expected result.

The results of testing our implementation with multiple learning rates and iteration counts are presented in Figure 1. Smaller and larger values than displayed have been tried out and have been discarded.

A short number of iterations is preferable (given the same accuracy). The best set of hyperparameters is thus a learning rate of ≈ 10 and an iteration count of ≈ 100 .

The cost given the number of iterations is displayed for 2 selected learning rates: For 10 (see Figure 2), which is near the optimal learning rate (which we know from Figure 1), and for a substantially lower learning rate of 0.1 (see 3). We can see that the learning rate has a dramatic impact on the number of learning iterations that is required for good predictions.

III. INTERPRETATION

There is only one layer of hidden units. The activations of these units for input data where unit i takes the value 1 and the rest takes the value 0 are thus identical to the weights and biases coming in from unit i .

²<https://python-poetry.org/docs>

³<https://python-poetry.org/docs/basic-usage>




Fig. 1: Accuracy given the learning rate and the number of iterations.




Fig. 3: For comparison: Cost given the number of iterations for a fixed learning rate of 0.1.




Fig. 2: Cost given the number of iterations for a fixed learning rate of 10. This is near the optimal learning rate.

We observe that the weights encode the 8 categories in a binary fashion – see Figure 4, and compare with Figure 5! This is the most efficient encoding possible, since $8 = 2^3$ categories have to be compressed into 3 units. It is positively surprising that the network so intelligently finds this representation.

We verify our hypothesis about the efficient encoding by training a network of shape $16-4-16$ with 16 different vectors in analogy to the 8-vector training set used before. The results can be seen in Figure 6: Indeed, we can again find the truth-table-like structure in the trained weights.

REFERENCES

- [1] A. Ng, “Sparse autoencoder,” *CS294A Lecture notes*, vol. 72, no. 2011, pp. 1–19, 2011.




Fig. 4: Visualization of the weights and biases (normalized by a sigmoid function) for the 8 input data. Connections between input units and hidden units (left) and between hidden units and output units (transposed, right). The former are identical to the activations of the hidden layer for the respective data. The order of the rows has been adjusted manually.

0	0	0	0
1	0	0	1
2	0	1	0
3	0	1	1
4	1	0	0
5	1	0	1
6	1	1	0
7	1	1	1

Fig. 5: Binary encodings of the first eight numbers, for comparison.




Fig. 6: Visualization of the weights and biases (normalized by a sigmoid function) for 16 input data. Connections between input units and hidden units. The rows have been ordered manually to match the truth-table for 4 bits. We use the same hyperparameters as above (a learning rate of 10), but extend the number of iterations from 100 to 200.